

Building Deep Neural Networks

with

Daniel L. Silver, Ph.D.

Andy McIntyre, Ph.D.


June 24, 2022

Outline

- Background
- Problem of vanishing gradient
 - Beyond two hidden layers
- Methods of overcoming the problem
 - Long patient training
 - Better activation functions
 - Pre-training
 - Constrained internal representation - weight-sharing
 - Better regularizers (sparsity, weight-decay, dropout, injected noise)
- Demo using many layers


Background – Why Deep Networks?

- Accuracy
 - Overcome limitations of shallow networks
 - Upper hidden layers learn higher order features = combinations of lower features
- Applications:
 - Image classification (ImageNet, medical analyses)
 - Voice recognition (Amazon Echo, Google Home)
 - Human competitive results (Game of GO)




Exploding or Vanishing (Unstable) Gradients

- What happens to the magnitude of the gradients as we backpropagate through many layers?
 - If the weights are big the gradients grow exponentially (explode)
 - If the weights are small, the gradients shrink exponentially (vanish)
- Typical feed-forward neural nets can cope with these exponential effects because of few hidden layers
- Deep Networks
 - Exploding gradient - clip the weights
 - Vanishing gradient - bigger problem
 - Why? Check out this video
- Details: <https://ayearofai.com/rohan-4-the-vanishing-gradient-problem-ec68f76ffb9b>




Vanishing Gradient

- Lowest layers of MLP do not get trained well
- Leads to very slow training - especially in lower layers
- Top couple layers can usually learn any task "pretty well"
- The error to lower layers drops quickly
- Lower layers never get the opportunity to use their capacity to improve results, they just create a random feature map
- Need a way for early layers to do effective work






Overcoming Vanishing/Exploding Gradient

- Better weight initialization
- Perform long patient training
 - small learning rates
- Use parallel processing - GPUs, etc.
- Normalize input to layers
- Use better activation functions
- Regularization through weight control



Better Activation Functions



- **Rectified Linear Units or Relu:** $f(x) = \text{Max}(0, x)$ - more efficient gradient propagation, derivative is 0 or constant, just fold into learning rate
- **Variations on Relu More efficient computation:** Only comparison, addition and multiplication.
 - Leaky ReLU $f(x) = x$ if $x > 0$ else ax , where $0 \leq a \leq 1$, so that derivative is not 0 and can do some learning for net < 0 (does not “die”).
 - Lots of other variations
- **Sparse activation:** For example, in a randomly initialized network, only about 50% of hidden units are activated (have output $\neq 0$)
- Learning in linear range easier for most learning models

Better Regularization of Free Parameters

- Weights (free parameters) make the ANN model
- Beyond the training examples we can control weights
 - Momentum
 - Weight-decay
 - Normalization
 - Sparsity
 - Dropout
 - Injected noise




Figure 1: Dropout Neural Net Model. **Left:** A standard neural net with 2 hidden layers. **Right:** An example of a thinned net produced by applying dropout to the network on the left. Crossed units have been dropped.

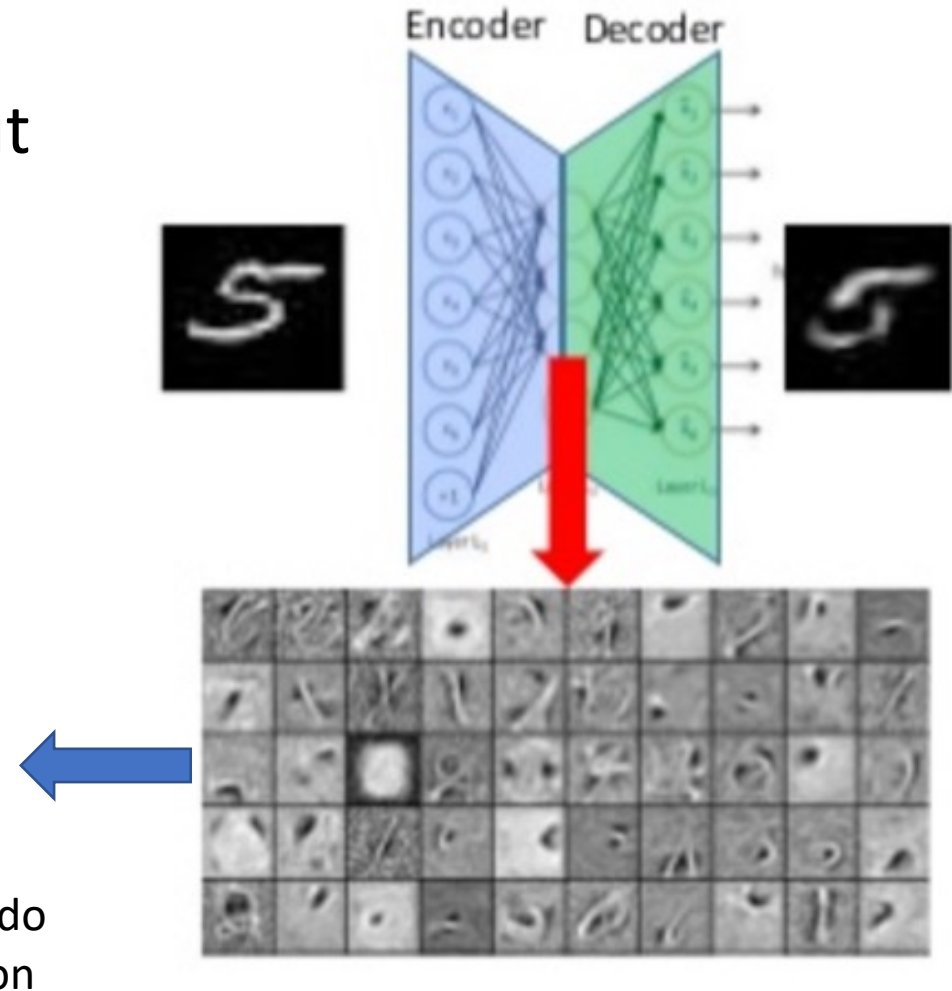
Pre-training of Deep Learning Architectures

- Geoff Hinton and Yoshua Bengio (2007+)
 - Learning Internal Representations
 - Layered networks of unsupervised auto-encoders efficiently develop hierarchies of features that capture regularities in their respective inputs
 - Auto-encoders
 - Restricted Boltzmann Machines
 - Can be used to develop models for families of tasks



Autoencoders using BP ANN

- Learn to predict their input at the output
- Learn an internal representation (encoding) of image
- Feature detectors of the input
 - Latent representation
 - Embedding space
- Can be used to pre-train a classifier




Deep Learning Architectures

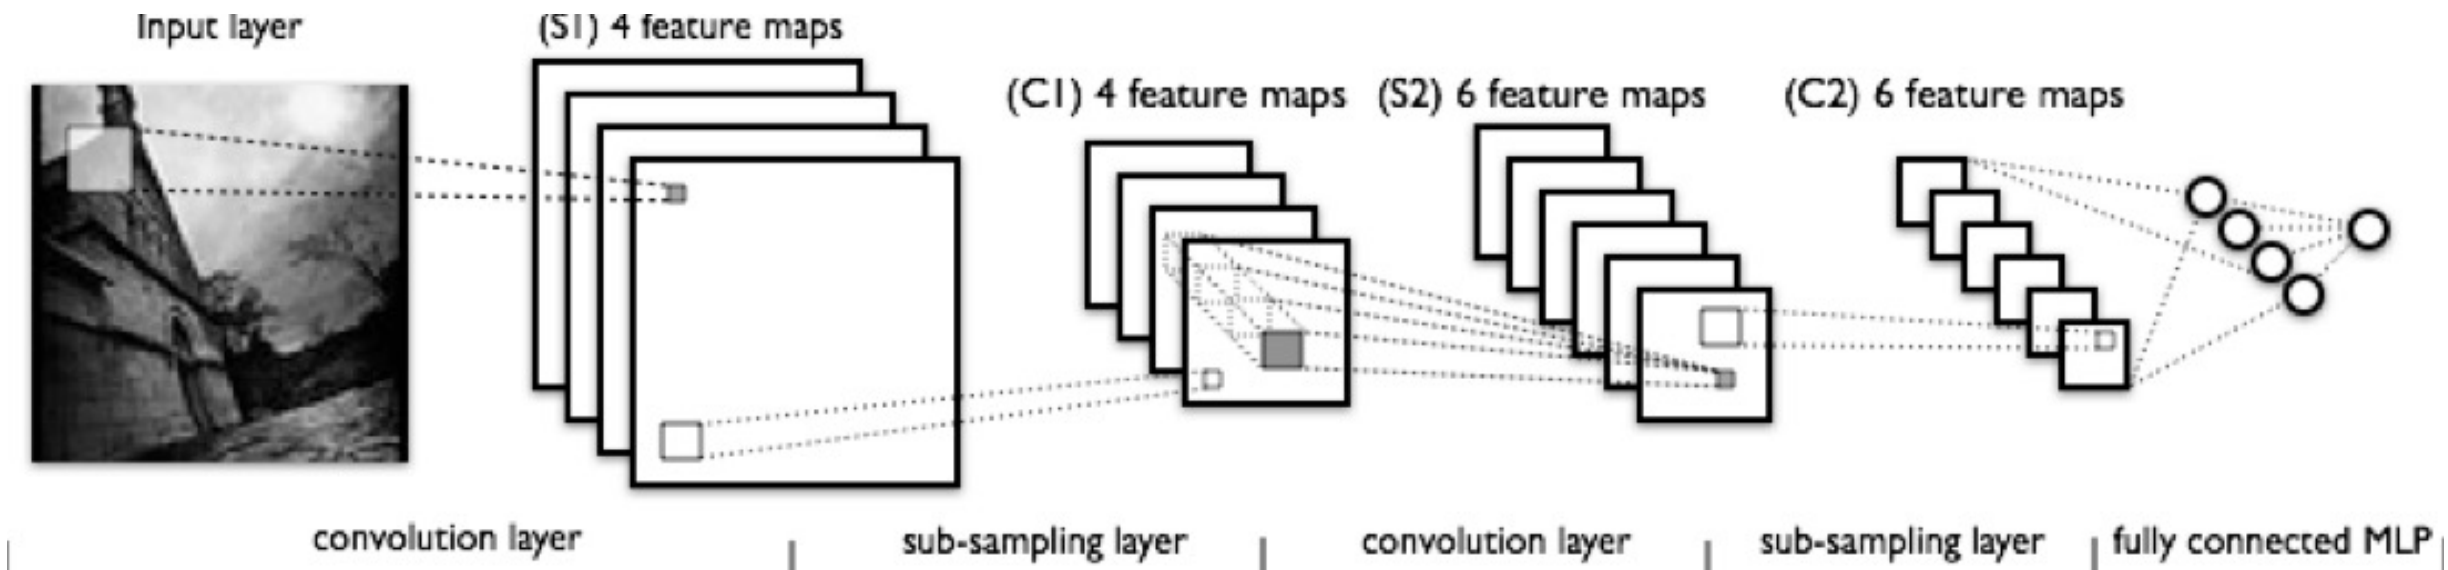
Andrew Ng's work on Deep Learning Networks (ICML-2012)

- Problem: Learn to recognize human faces, cats, etc from unlabeled data
- Dataset of 10 million images; each image has 200x200 pixels
- 9-layered locally connected sparse autoencoder neural network (1B connections)
- Parallel algorithm; 1,000 machines (16,000 cores) for three days

Building High-level Features Using Large Scale Unsupervised Learning
Quoc V. Le, Marc' Aurelio Ranzato, Rajat Monga, Matthieu Devin, Kai Chen, Greg S. Corrado, Jeffrey Dean, and Andrew Y. Ng
ICML 2012: 29th International Conference on Machine Learning, Edinburgh, Scotland, June, 2012.



Constrained Internal Representation



Weight sharing based on knowledge of the input space

- Reduces overall complexity of model
- One of key hints for using deep learning architectures

TUTORIAL 4

- Demonstrate how the addition of hidden layers makes the development / training of a BP network more challenging
- Try the python code **`tf.keras_mnist_deep.ipynb`**
- For more, check out the following tutorial:
<https://www.tensorflow.org/tutorials/keras/classification>