


Convolution Neural Networks

with

Daniel L. Silver, Ph.D.

Andy McIntyre, Ph.D.

June 24, 2022



**ACADIA INSTITUTE
FOR DATA ANALYTICS**
digital harvest

[Slides based on originals by Yann LeCun]



Convolution Neural Networks

- Fundamentally CNNs take advantage of knowledge that we have about the problem's input space
- We know that pixels in an image that are adjacent to each other are related – similar color and brightness
- We can use this background knowledge as a source of inductive bias to help develop better NN models



Receptive Fields

- The **receptive field** of an individual sensory neuron is the particular region of the sensory space (e.g. the retina) in which a stimulus will trigger the firing of that neuron.
 - It is a feature detector (image line orientation, sound frequency)
- How do we design “proper” receptive fields for the input neurons?
- Consider a task with image inputs
 - Receptive fields should provide expressive features from the raw input
 - How would you design the receptive fields for this problem?



Solution 1


■ A fully connected layer and network:

□ Example:

- 100x100 images
- 1000 units in 1st hidden layer

□ Problems:

- 10^7 edges!
- Local spatial correlations lost!
- Variable sized inputs.



Slide Credit: Marc'Aurelio Ranzato

Solution 2

■ A locally connected layer:


□ Example:

- 100x100 images
- 1000 units in the input
- Filter size: 10x10

□ Local correlations preserved!

□ Problems:

- 10^5 edges
- This parameterization is good when input image is registered (e.g., face recognition).
- Variable sized inputs, again.



Variation of Objects In Images

- Fixed objects can vary in terms of:
 - Translation (location)
 - Rotation
 - Scale
- Handwritten digits can vary due to nuances of pen stroke




Better Solution - Convolution

■ A solution:

- ❑ **Filters** to capture different patterns in the input space.
 - **Share** parameters across different locations or rotations
 - **Convolutions** with learned filters
- ❑ Filters will be **learned** during training.
- ❑ The issue of variable-sized inputs will be resolved with a **pooling** layer.

So what is a convolution?



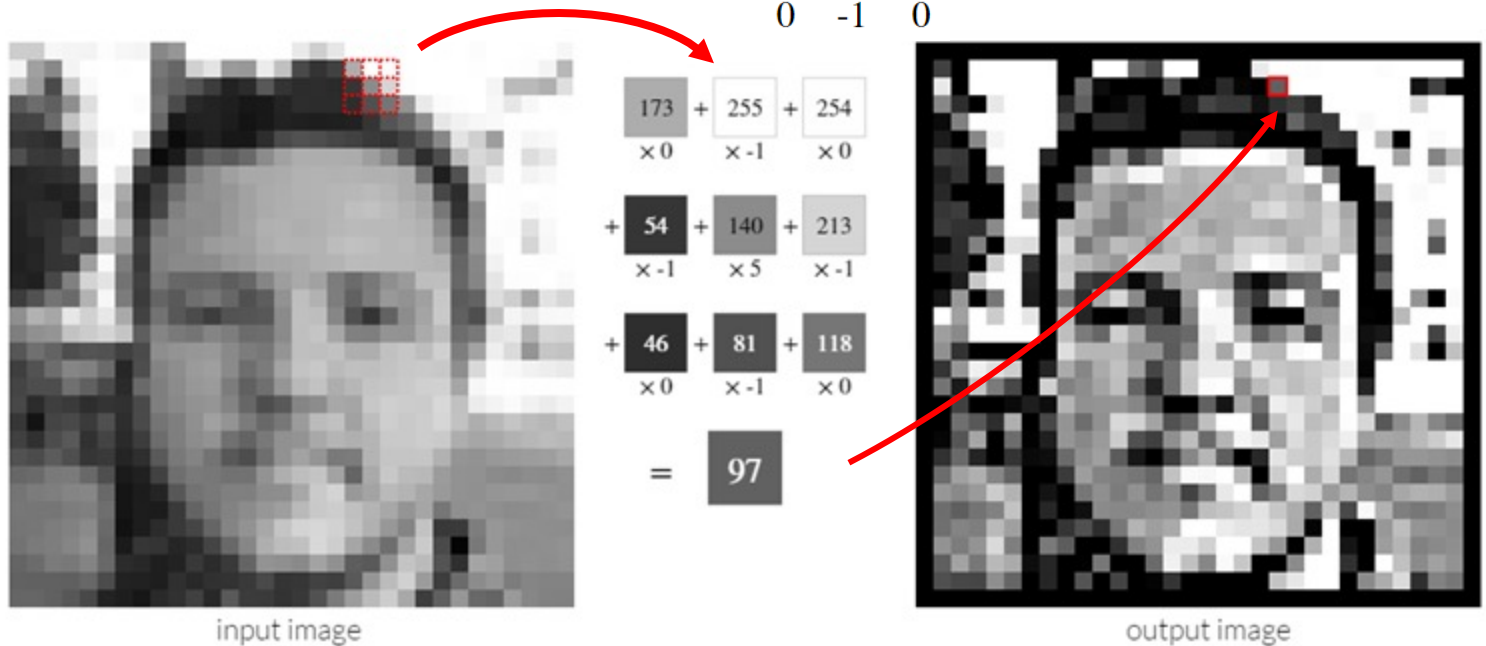
Slide Credit: Marc'Aurelio Ranzato

Convolution Operator

- Convolution in two dimension:


- Takes two functions and produces another function
- 2D: Take one matrix and slide it over the other matrix
- Example: Sharpen kernel:

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$



Try other kernels: <http://setosa.io/ev/image-kernels/>

Convolution Operator



Example of Blurring
Filter

0	0	0	0	0
0	1	1	1	0
0	1	1	1	0
0	1	1	1	0
0	0	0	0	0


Example of Edge
Detector

	0	1	0	
	1	-4	1	
	0	1	0	

Convolution Operator

- Convolution in two dimension:
 - The same idea: flip one matrix and slide it on the other matrix




6 x 6 x 1 image



Slide Credit: Marc'Aurelio Ranzato

Convolutional Layer

- The convolution of the **input (vector/matrix)** with weights **(vector/matrix)** results in a **response vector/matrix**.
- We can have **multiple (weight-based) filters** in each convolutional layer, each producing an output **feature map**
- If it is an intermediate layer, it can have **multiple inputs**




Feature Map



One can add nonlinearity
at the output of
convolutional layer

Pooling Layer

- How to handle variable sized inputs?
 - A layer which reduces inputs of different size, to a fixed size.
 - **Pooling (also called subsampling)**



Slide Credit: Marc'Aurelio Ranzato

Pooling Layer

■ How to handle variable sized inputs?

- A layer which reduces inputs of different size, to a fixed size.
- **Pooling (also called subsampling)**
- Different variations

- Max pooling

$$h_i[n] = \max_{i \in N(n)} \tilde{h}[i]$$


- Average pooling

$$h_i[n] = \frac{1}{n} \sum_{i \in N(n)} \tilde{h}[i]$$

- L2-pooling


$$h_i[n] = \frac{1}{n} \sqrt{\sum_{i \in N(n)} \tilde{h}^2[i]}$$

- etc




Convolutional Nets


- One stage structure:



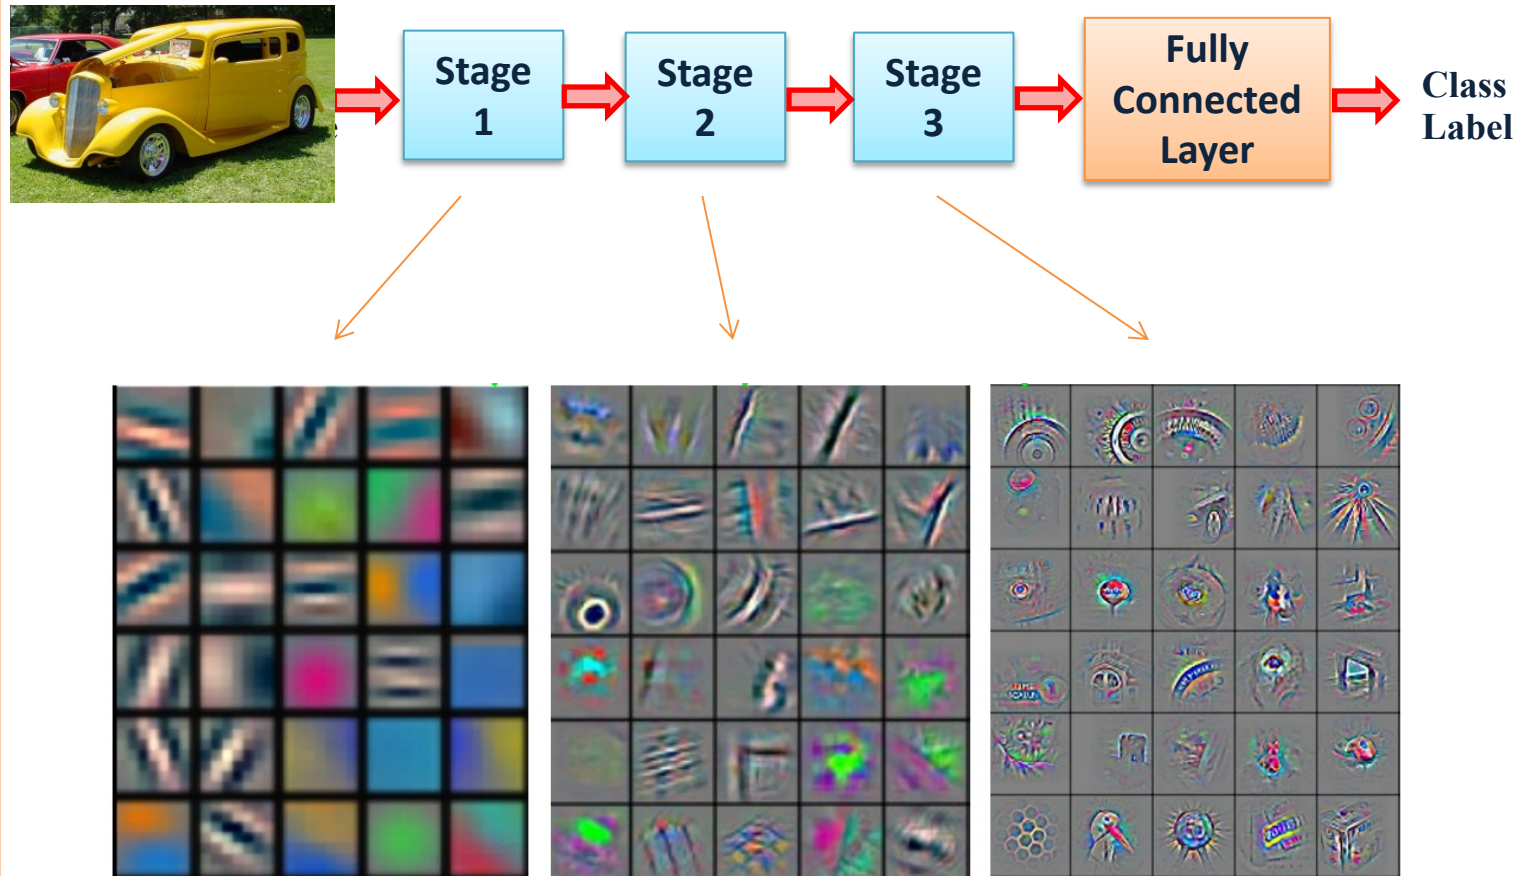
- Whole system:



An example
system (LeNet):



Convolutional Nets




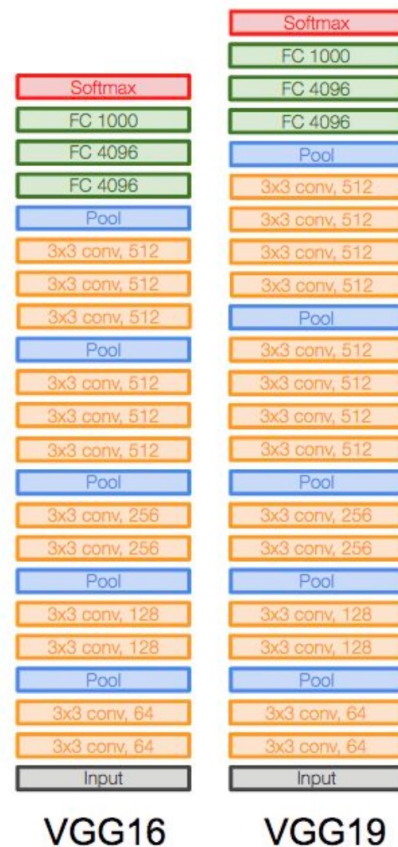


An example
system :

Feature visualization of convolutional net trained on ImageNet from [Zeiler & Fergus 2013]

Convolutional Nets

Advances in CNNs

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners

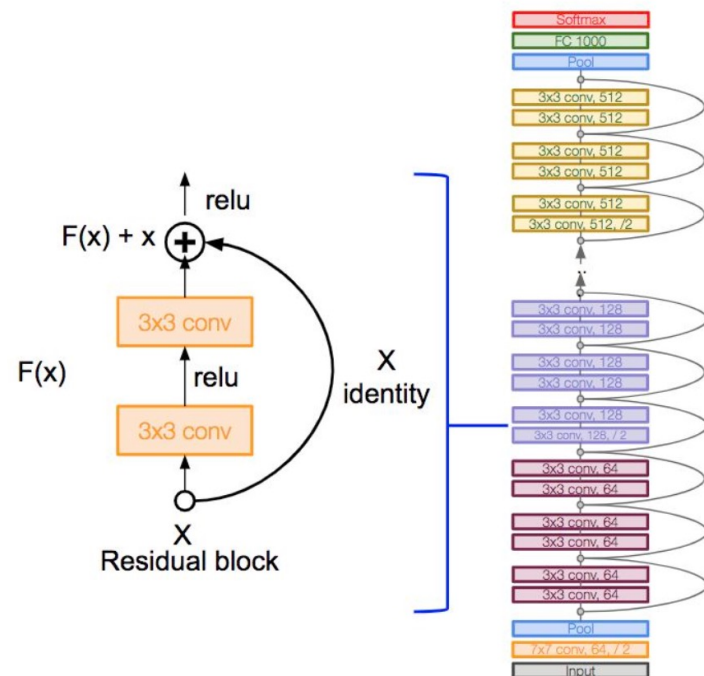




[Source: Fei-Fei Li, Ranjay Krisna, Danfei Xu / Stanford]

Convolutional Nets

Advances in CNNs

ImageNet Large Scale Visual Recognition Challenge (ILSVRC) winners



ResNet

[Source: Fei-Fei Li, Ranjay Krisna, Danfei Xu / Stanford]

Practical Tips

- Before large scale experiments, test on a small subset of the data and check the error should go to zero.
 - Overfitting on small training
- Visualize features (feature maps need to be uncorrelated) and have high variance
- Bad training: many hidden units ignore the input and/or exhibit strong correlations.

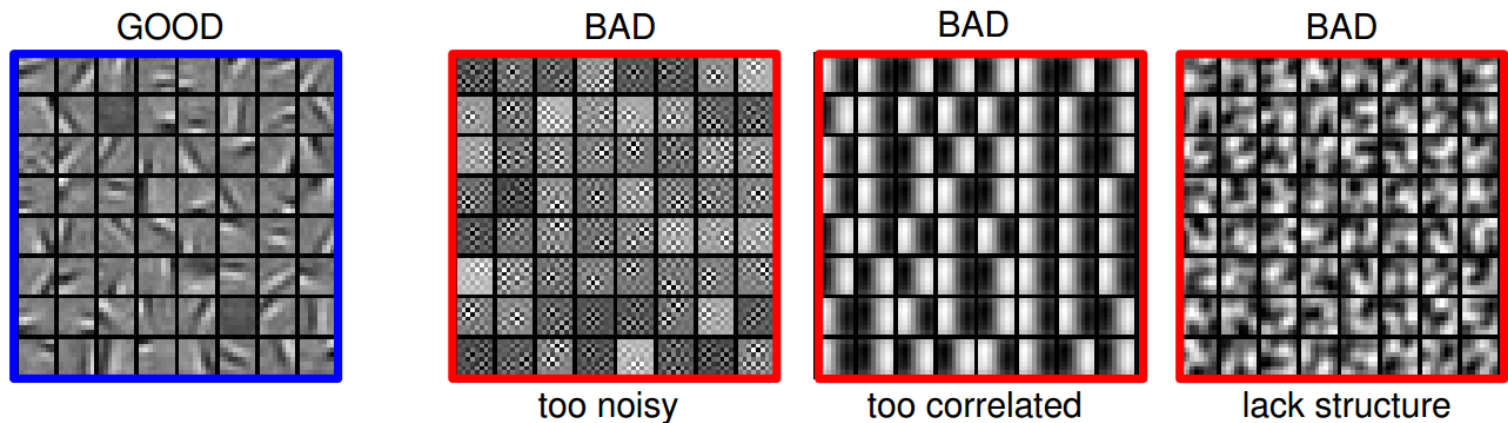


Figure Credit: Marc'Aurelio Ranzato

Dropout typically useful




Figure 1: Dropout Neural Net Model. **Left:** A standard neural net with 2 hidden layers. **Right:** An example of a thinned net produced by applying dropout to the network on the left. Crossed units have been dropped.

TUTORIAL 4

- Develop and train a CNN network using Keras and Tensorflow (keras_mnist_cnn_val.ipynb)

References

- Demo: <https://www.cs.cmu.edu/~aharley/vis/conv/>
- Further Reading:
 - <https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/>
 - <https://adeshpande3.github.io/A-Beginner's-Guide-To-Understanding-Convolutional-Neural-Networks-Part-2/>